




INFERENCE API EVALUATION GUIDE

SUPERAI ENGINEER #SS6 – UPDATE 20260603_1915



EVELUTION ARCHITECTURE



API SPECIFICATION FOR EVALUATION – OCR

API Specification for Evaluation	
Endpoint URL	<u>http://swarm-manager.modelharbor.com:YOUR_PORT/ocr</u>
Method	POST *Or your setting URL*
Header	Content-Type: application/json
Request Body	<pre>{ "id": "QUESTION_ID", "header": "BASE64_ENCODED", "transaction": ["BASE64_ENCODED"] }</pre>
Response Body	<pre>{ "id": "QUESTION_ID", "answer": {"FOLLOWING KAGGLE FORMAT"} }</pre>

API SPECIFICATION FOR EVALUATION – THAI LLM

API Specification for Evaluation

Endpoint URL	<u>http://swarm-manager.modelharbor.com:YOUR_PORT/agent/thai1lm</u>
--------------	---

Method	POST *Or your setting URL*
--------	-----------------------------------

Header	Content-Type: application/json
--------	--------------------------------

Request Body	<pre>{ "question": "ข้อความคำถาม" }</pre>
--------------	---

This link to your "Audit file"

Format:

YOUR_GENERTION_ID_FOR_EACH_QUESTION.txt

Response Body	<pre>{ "id": "YOUR_GENERTION_ID_FOR_EACH_QUESTION", "answer": "ข้อความคำตอบ", "total_output_token_count": int }</pre>
---------------	---

API SPECIFICATION FOR EVALUATION – LOCAL LLM

API Specification for Evaluation

Endpoint URL	<u>http://swarm-manager.modelharbor.com:YOUR_PORT/agent/local</u>
--------------	---

Method	POST *Or your setting URL*
--------	-----------------------------------

Header	Content-Type: application/json
--------	--------------------------------

Request Body	<pre>{ "question": "ข้อความคำถาม" }</pre>
--------------	---

This link to your "Audit file"

Format:

YOUR_GENERTION_ID_FOR_EACH_QUESTION.txt

Response Body	<pre>{ "id": "YOUR_GENERTION_ID_FOR_EACH_QUESTION", "answer": "ข้อความคำตอบ", "total_output_token_count": int }</pre>
---------------	---

EVALUATION PARAMETER FOR API APP PERFORMANCE

Jupyter

<https://canva.link/ny3ag44b24h8y7k>

Proxy Frontend

[http://swarm-manager.modelharbor.com:50235/hub/share?token=sharetoken&next=/user/\[USERNAME\]/proxy/\[PORT\]/](http://swarm-manager.modelharbor.com:50235/hub/share?token=sharetoken&next=/user/[USERNAME]/proxy/[PORT]/)

Proxy Backend API

[http://swarm-manager.modelharbor.com:50235/hub/public-proxy/\[USERNAME\]/\[PORT\]?_token=sharetoken](http://swarm-manager.modelharbor.com:50235/hub/public-proxy/[USERNAME]/[PORT]?_token=sharetoken)

Ubuntu (OS)

Port Mappings

Port 22

51090 → 22/tcp

Port 9000

50235 → 9000/tcp

Loadtest Engine

System host port

In-Container port

<http://swarm-manager.modelharbor.com:50235/>

EVALUATION PARAMETER FOR API APP PERFROMANCE

- Stress Parameter
 - Number of concurrent users: 10
 - User growth per second: 2
- Measure Parameter
 - No of reqs
 - No of fails
 - Avg/Min/Max/Median time (second)
 - Req per second
 - Failures per second
- Tool
 - locust.io

EVALUATION PARAMETER FOR API SERVER RESOURCE UTILIZATION

- Stress Parameter
 - Number of concurrent users: 10
 - User growth per second: 2
- Measure Parameter (Every 1 second)
 - GPU-Memory Usage
 - GPU-Utilization
- Tools
 - nvidia-smi

EVALUATION PARAMETER FOR API ANSWER ACCURACY SCORE

- Blind input dataset via API Request
- Measure Parameter
 - Accuracy score of answer

REMARK

- Image is PNG
- Start Testing
 - Thu 04 June 2569
 - 07:30 AM ← OCR Testing
 - 08:30 AM ← ThalLLM Testing → Easy=5, Medium=9, Hard=1
 - 09:00 AM ← LocalLLM Testing
- After 07:30 AM, Any remote activities to following instruct from “Committee”. ****Any apart of this are prohibit.****