


Using AI modules on LANTA

Yutthana Wongnongwa, Ph.D.

NSTDA Supercomputer Center (ThaiSC)
National Electronics and Computer Technology Center (NECTEC)
National Science and Technology Development Agency (NSTDA)



- **Basic knowledge of LANTA**
 - **System overview**
 - **Storage quota**
 - **Basic Lmod commands**
 - **Overview of Slurm job script**
 - **Checking job status**
 - **Cancel job**
 - **Check your billing account**
 - **File and folder permissions**
- **Using Miniforge3 module on LANTA**
- **Example of running Jupyter Notebook**
- **Example of Multi-GPUs/Multi-nodes training**
- **Using cray-python module on LANTA**
- **Using Apptainer module on LANTA**

Basic knowledge of LANTA

How to Login to LANTA

Windows

- ✓ Login to LANTA with **MobaXterm**
<https://youtu.be/NW40sB5W1dc>

- ✓ (Optional) Login to LANTA with **VSCode**
<https://youtu.be/4pzezKQZO3g>

MacOS

- ✓ Login to LANTA with **Terminal**
<https://youtu.be/kMTRAw-QHro>

- ✓ (Optional) Login to LANTA with **Termius**
<https://youtu.be/INdVbsNOS6c>

- ✓ (Optional) Login to LANTA with **VSCode**
<https://youtu.be/icueMyxhTZQ>

LANTA frontend policy

lanta.nstda.or.th

- ✓ Login to LANTA
- ✓ Use general commands
- ✓ Submit Job
- ✓ Internet Speed: **up to 800 Mbps**

Note: Do not use for file transfer and run heavy process (>90%)

transfer.lanta.nstda.or.th

- ✓ File transfer
- ✓ Create conda environment
- ✓ Download container file
- ✓ Internet Speed: **up to 20 Gbps**

Note: Do not run heavy process (>90%)

System overview

Memory Nodes

1280 Cores

10x Compute nodes

- **128 cores/node**
2 x AMD EPYC™ 7713 “Milan”
- 4 TB DRAM (ECC)

Compute Nodes

20,480 Cores

160x Compute nodes

- **128 cores/node**
2 x AMD EPYC™ 7713 “Milan”
- 256 GB DRAM (ECC)


GPU Nodes 704 GPU

176x GPU nodes


- **64 cores/node**
1 x AMD EPYC™ 7713 “Milan”
- 512 GB DRAM
- **4x Nvidia A100 40GB (HGX)**
- 2.4 TB Aggregated Bandwidth

High Performance Storage And Network

- 200 Gbps
- 10 PB HPC storage
- Connected all compute and storage via dragonfly network



Check LANTA status



Execute “**sinfo**” command to check LANTA status

```
PARTITION    AVAIL  TIMELIMIT  NODES  STATE NODELIST
compute*     up  5-00:00:00    46   mix lanta-c-[004-006,010,014-016,020,
compute*     up  5-00:00:00    24  alloc lanta-c-[001-003,007-009,011-013,
compute*     up  5-00:00:00    88   idle lanta-c-[022-064,114-158]
compute-devel up    2:00:00     2   idle lanta-c-[159-160]
gpu          up  5-00:00:00     5   mix lanta-g-[001-002,005-007]
gpu          up  5-00:00:00    10  alloc lanta-g-[003-004,009-016]
gpu          up  5-00:00:00   159   idle lanta-g-[008,017-174]
gpu-devel    up    2:00:00     2   idle lanta-g-[175-176]
memory       up  5-00:00:00     4   mix lanta-m-[001-004]
memory       up  5-00:00:00     5  alloc lanta-m-[005-009]
memory       up  5-00:00:00     1   idle lanta-m-010
```

State	Detail
idle	The whole machine is available to use.
alloc	The whole machine is currently in use.
mix	Some resources of the machine are available to use, while some are already in use.
down	The machine is NOT available for use.
drain	The machine is NOT available for further use, due to some minor issues with the system.

LANTA storage quota

Path	Storage space	Inode/count	Permission
/home/username	100 GB	600,000	For each user account
/project/zz991000-zdeva/zz991000	200 TB	2,000,000,000	Share among all project members
/project/zz991000-zdeva/zz991xxx	200 TB	2,000,000,000	Share among all project members



*** Important***

To check remaining quota,
execute “myquota” command

	Block Quota Limits		Inode Quota Limits	
	Used	Size	IUsed	Inodes
/home/hpcuser	31.57G	100G	172925	600000
/project/lt200199-espre	61.81G	5T	192973	50000000

Basic Lmod commands

Command	Subcommand	<Arg (0) ...>	= action
module	overview		= list all available modules by short names and count of versions
module	help	<i>name</i>	= print short detail of <i>name</i>
module	spider	<i>name</i>	= print detailed information of <i>name</i> (slow)
module	avail	<i>name</i>	= list all available versions of <i>name</i> ; if none specified, then list all
module	list		= list all modules currently in use
module	load	<i>name/version</i>	= load <i>name/version</i> module
module	unload	<i>name/version</i>	= unload <i>name/version</i> module
module	swap	<i>old new</i>	= replace <i>old</i> with <i>new</i>
module	purge		= unload all modules
module	reset		= restore to the <u>default</u> set of modules (LANTA: PrgEnv-cray)

*** Tips ***

Use “ml” instead “module”
and “module list”

Overview of Slurm job script

Slurm job scripts enable the execution of software on hardware resources for a specified duration

Example: job script

```
#!/bin/bash
#SBATCH --partition=partition_name
#SBATCH --nodes=num_nodes
#SBATCH --ntasks-per-node=ntask/node
#SBATCH --gpus-per-node=ngpu/node
#SBATCH --time=D-HH:MM:SS
#SBATCH --job-name=jobname
#SBATCH --account=ltxxxxxx
```

```
module purge
module load libs-or-progs

srun my_hpc_application
```

1) #SBATCH directive/header

- State required resources
- Various possible approaches/combinations, see [sbatch](#)

2) Instruction/Script

- Run on the first scheduled node
 - Load/Set software environment
 - Execute instructions
- Use [srun](#) to launch software processes across nodes
- By default, **srun inherits most options from sbatch**; since all allocated resources should be utilized.

Check job status

Execute “**myqueue**” command to check your job status

```

username@x3002c0s7b0n0:~> myqueue
JOBID      PARTITION  NAME      USER      ST          TIME      NODES      NODELIST (REASON)
151602     compute   test4     username   PD          0:00       1          (Resources)
151600     compute   test3     username   PD          0:00       1          (Priority)
150500     compute   test3     username   PD          0:00       1          (QOSGrpBillingMinutes)
140215     compute   test2     username   R    03:25:08   1          lanta-c-055
149496     compute   test1     username   R 2-02:11:27  1          lanta-g-042
  
```

(REASON)

- (Resources) = Not enough resources for your job. However, **your job is the 1st in the queue.**
 - (Priority) = Not enough resources for your job. **Your job is waiting in the queue.**
 - (QOSGrpBillingMinutes) = Remaining SHr of the project is NOT enough for the job to run until reaching its runtime limit.
- Check by executing “sbalance” --> try reducing the runtime lime (-t) or top up more SHr



Cancel job

Execute “**scancel [JOBID]**” command to cancel job

```
username@x3002c0s7b0n0:~> myqueue
```

JOBID	PARTITION	NAME	USER	ST	TIME	NODES	NODELIST (REASON)
151602	compute	test4	username	PD	0:00	1	(Resources)
151600	compute	test3	username	PD	0:00	1	(Priority)
150500	compute	test3	username	PD	0:00	1	(QOSGrpBillingMinutes)
140215	compute	test2	username	R	03:25:08	1	lanta-c-055
149496	compute	test1	username	R	2-02:11:27	1	lanta-g-042

```
username@x3002c0s7b0n0:~> scancel 140215
```

```
username@x3002c0s7b0n0:~> myqueue
```

JOBID	PARTITION	NAME	USER	ST	TIME	NODES	NODELIST (REASON)
151602	compute	test4	username	PD	0:00	1	(Resources)
151600	compute	test3	username	PD	0:00	1	(Priority)
150500	compute	test3	username	PD	0:00	1	(QOSGrpBillingMinutes)
149496	compute	test1	username	R	2-02:11:27	1	lanta-g-042

Check your billing account

Execute “**sbalance**” command to check SHr of your billing accounts

```
username@x3002c0s7b0n0:~> sbalance
```

Account	Descr	Remaining(%)	Allocation(SHr)	Used(SHr)	Remaining(SHr)
1t200173	cperft	37.10%	4281.00	2692.94	1588.06
1txxxxxx	yyyyy	40.16%	221875.00	132764.58	89110.42

- Use “**sbalance --detail**” to display (approximated) usage per users
- See “**sbalance --help**” for more available options

Billing of allocated resource

Whole-node allocation

	CPU cores	GPU cards	RAM* (GiB)	Cost per hour
lanta-c-xxx	128	-	up to 250G	1 SHr
lanta-g-xxx	64	4	up to 502G	3 SHr
lanta-m-xxx	128	-	up to 4018G	4 SHr

Standard allocatable pack of resources

	CPU cores	GPU cards	RAM* (MiB)	Cost per hour
lanta-c-xxx	1	-	1,900M	~0.0078 SHr
lanta-g-xxx	16	1	124,000M	0.75 SHr
lanta-m-xxx	1	-	32,000M	0.03125 SHr

RAM* = Free RAM, readily for software to use
 1 GiB = 1024 MiB

*** Warning ***

This implies that

- 32 CPU + 2 GPU,
- 2 CPU + 2 GPU,
- 32 CPU + 1 GPU

equally cost $2 \times 0.75 = 1.5$ SHr per hour, because 2 packs (half GPU node) are retained for your job.

Note: 1) For lanta-g-xxx, every extra CPU core + 8 GiB RAM will cost 0.046875 SHr per hour, and then every further ~1,365 MiB RAM will cost ~0.0078 SHr per hour on top.
 2) For lanta-m-xxx, every extra 8,000 MiB RAM will cost ~0.0078 SHr per hour on top.

File and folder permissions

Execute “**ll**” command to check your file and folder permissions

```
username@x3002c0s7b0n0:~> ll
drwxr-xr-x  11 username lt200xxx  4096 Apr 18 16:14 Your_Directory
```

Execute “**chmod g+w**” command to change your file and folder permissions

```
username@x3002c0s7b0n0:~> chmod g+w Your_Directory
username@x3002c0s7b0n0:~> ll
drwxrwxr-x  11 username lt200xxx  4096 Apr 18 16:14 Your_Directory
```

- **u** - The file owner.
- **g** - The users who are members of the group.
- **o** - All other users.

Permission	Character	Meaning on File
Read	-	The file is not readable. You cannot view the file contents.
	r	The file is readable.
Write	-	The file cannot be changed or modified.
	w	The file can be changed or modified.
Execute	-	The file cannot be executed.
	x	The file can be executed.

FAQ: Incorrect storage quota

- “myquota” will **double count** files in /project/ path, if their group is your username.
- This usually happens when `mv` command is used to move files between /project/ and /home/

```
skanoksi@x3002c0s9b0n0:/project/thaisc/skanoksi/WorkSpace/tutorial> ls -lR
.:
total 8
drwxr-sr-x 2 skanoksi thaisc 4096 dirA
drwxr-xr-x 2 skanoksi thaisc 4096 dirB

./dirA:
total 13778020
-rw-r--r-- 1 skanoksi thaisc 14108684400 largefile

./dirB:
total 13762564
-rw-r--r-- 1 skanoksi skanoksi 14108684400 largefile
```

Count in
/project/ (5,120 GB)

Count in both
/project/ (5,120 GB) and
/home (100GB) !!

To fix this issue, use

1. **`chgrp -Rf ltxxxxxx your-directory`**
= to change group of all files inside your-directory
2. **`find your-directory -type d -exec chmod g+s {} \;`**
= to add setuid bit (+s) to group permission (g) of all directories (-type d) inside your-directory

Using Miniforge3 module on LANTA

➤ Load the Mamba module and activate your environment

```
username@lanta:~> ml Miniforge3/25.3.0-3
username@lanta:~> conda activate environment-name
(environment-name) username@lanta:~>
```

➤ View a list of your environments

```
username@lanta:~> conda env list
# conda environments:
#
base                /lustrefs/disk/modules/easybuild/software/Miniforge3/25.3.0-3
netcdf-py39         /lustrefs/disk/modules/easybuild/software/Miniforge3/25.3.0-3/envs/netcdf-py39
pytorch-2.8.0       /lustrefs/disk/modules/easybuild/software/Miniforge3/25.3.0-3/envs/pytorch-2.8.0
tensorflow-2.20.0   /lustrefs/disk/modules/easybuild/software/Miniforge3/25.3.0-3/envs/tensorflow-2.20.0
```

➤ Deactivate your environment

```
(environment-name) username@lanta:~> conda deactivate  
username@lanta:~>
```

➤ Unload the Miniforge3 module

```
username@lanta:~> module unload Miniforge3
```

➤ Unload all currently loaded modules

```
username@lanta:~> module purge
```

Note: Before you use the **module unload Miniforge3** or **module purge** command, you must **deactivate your environment** with the **conda deactivate** command.

➤ Create the environment

```
username@lanta:~> conda create -n myenv  
.  
.  
proceed ([y]/n)?
```

➤ Create the environment with a specific version of packages

```
username@lanta:~> conda create -n myenv python=3.9
```

```
username@lanta:~> conda create -n myenv python=3.9 numpy=1.23.5
```

➤ Activate your environment

```
username@lanta:~> conda activate myenv
```

➤ Specify a location for the environment

```
username@lanta:~> conda create --prefix /your project directory/envs
```

➤ Specify a location for the environment with a specific version of a packages

```
username@lanta:~> conda create --prefix /your project directory/envs python=3.9 numpy=1.23.5
```

➤ Activate your environment

```
username@lanta:~> conda activate /your project directory/envs
```

➤ A simple requirements.yml file

```
name: myenv
dependencies:
  - python=3.9
  - numpy=1.23.5
  - pandas
```

➤ Create the environment from the requirements.yml file in the user's home

```
username@lanta:~> conda env create -f requirements.yml
```

➤ Create the environment from the requirements.yml file in the project's home

```
username@lanta:~> conda env create -f requirements.yml --prefix /your project directory/envs
```

➤ Create the R environment in the user's home

```
username@lanta:~> conda create -n r_env r-essentials r-base
.
.
proceed ([y]/n)?
```

➤ Create the R environment in the project's home

```
username@lanta:~> conda create --prefix /your project directory/r_env
r-essentials r-base
.
.
proceed ([y]/n)?
```

➤ Activate your environment

```
username@lanta:~> conda activate r_env
username@lanta:~> conda activate /your project directory/r_env
```

➤ Running on Compute node

```
#!/bin/bash
#SBATCH -p compute
#SBATCH -N 1 -c 128
#SBATCH --ntasks-per-node=1
#SBATCH -t 120:00:00
#SBATCH -A tb901xxx
#SBATCH -J JOBNAME

module load Miniforge3/25.3.0-3
conda activate tensorflow-2.20.0

python3 file.py
```

Specify partition [Compute/Memory/GPU]
Specify number of nodes and processors per task
Specify tasks per node
Specify maximum time limit (hour: minute: second)
Specify project name
Specify job name

Load the Mamba module
Activate your environment

Run your program or executable code

Note: Full node: -c 128, Half node: -c 64, ¼ node: -c 32

➤ Running on GPU node

```
#!/bin/bash
#SBATCH -p gpu
#SBATCH -N 1 -c 16
#SBATCH --gpus-per-node=4
#SBATCH --ntasks-per-node=4
#SBATCH -t 120:00:00
#SBATCH -A tb901xxx
#SBATCH -J JOBNAME

module load Miniforge3/25.3.0-3
conda activate tensorflow-2.20.0

python3 file.py
```

Specify partition [Compute/Memory/GPU]
Specify number of nodes and processors per task
Specify number of GPU per task
Specify tasks per node
Specify maximum time limit (hour: minute: second)
Specify project name
Specify job name

Load the Mamba module
Activate your environment

Run your program or executable code

Note: 1 GPU card: --gpus-per-node=1, 2 GPU cards: --gpus-per-node=2, 4 GPU cards: --gpus-per-node=4

Example of running Jupyter Notebook

Using the example to run Jupyter Notebook

```
username@lanta:~> cd /project/your_project
username@lanta:~> cp -r /project/common/Miniforge3/ .
username@lanta:~> cd Mamba/
username@lanta:~> ls
DDP_1N  DDP_2N  Finetune_BERT_1N  Finetune_BERT_2N  Jupyter_GPU
```

Example of running Jupyter Notebook

```
port=$(shuf -i 6000-9999 -n 1)
USER=$(whoami)
node=$(hostname -s)
```

```
# jupyter notebook instructions to the output file
echo -e "
```

```
Jupyter server is running on: $(hostname)
Job starts at: $(date)
```

Copy/Paste the following command into your local terminal

```
-----
ssh -L $port:$node:$port $USER@lanta.nstda.or.th -i id_rsa
-----
```

Open a browser on your local machine with the following address

```
-----
http://localhost:${port}/?token=XXXXXXXXX (see your token below)
-----
```

"

```
# start a cluster instance and launch the jupyter server
```

```
unset XDG_RUNTIME_DIR
if [ "$SLURM_JOBTMP" != "" ]; then
    export XDG_RUNTIME_DIR=$SLURM_JOBTMP
fi
```

```
jupyter notebook --no-browser --port $port --notebook-dir=$(pwd) --ip=$node
```


➤ There are 3 steps to run Jupyter Notebook on LANTA

1. Submit your job and read your slurm-xxxxx.out

```
username@lanta:~> sbatch script.sh  
username@lanta:~> cat slurm-xxxxx.out
```

2. Copy/Paste the following command into your local terminal

```
ssh -L 7637:x1000c2s0b0n0:7637 username@lanta.nstda.or.th -i id_rsa
```



```
Command Prompt - ssh -L 81  x + v  
Microsoft Windows [Version 10.0.22621.1265]  
(c) Microsoft Corporation. All rights reserved.  
  
C:\Users\006501>ssh -L 8147:x1000c2s0b0n0:8147 yutthana@lanta.nstda.or.th -i id_rsa  
Last login: Thu Feb 23 22:24:45 2023 from 10.223.1.138  
yutthana@x3002c0s7b0n0:~> |
```

➤ There are 3 steps to run Jupyter Notebook on LANTA

3. Open a browser on your local machine with the following address (Final line in .out)

`http://127.0.0.1:8147/?token=9720bf83310e36e6d6adb205ac3d2dc01c0518b73fb8421e`

Home Page - Select or create a notebook

127.0.0.1:8147/tree

jupyter

Quit Logout

Files Running Clusters

Select items to perform actions on them.

Upload New

	Name	Last Modified	File size
☐	0		
☐	mnist.ipynb	42 minutes ago	2.32 MB
☐	Jupyter_GPU_PT.sub	11 days ago	1.49 kB
☐	Jupyter_GPU_TF.sub	an hour ago	1.65 kB
☐	mnist.py	21 days ago	2.63 kB
☐	mnist_data.npy	21 days ago	55 MB
☐	setup.py	21 days ago	440 B
☐	slurm-28973.out	4 minutes ago	10.4 kB

Example of Multi-GPUs/Multi-nodes training

Using the example to run Multi-GPUs/Multi-nodes training

```
username@lanta:~> cd /project/your_project
username@lanta:~> cp -r /project/common/Miniforge3/ .
username@lanta:~> cd Mamba/
username@lanta:~> ls
DDP_1N  DDP_2N  Finetune_BERT_1N  Finetune_BERT_2N  Jupyter_GPU
```


Example of Multi-GPUs/Multi-nodes training

```
#!/bin/bash
#SBATCH -p gpu                # Specify partition [Compute/Memory/GPU]
#SBATCH -N 1 -c 16            # Specify number of nodes and processors per task
#SBATCH --ntasks-per-node=4   # Specify number of tasks per node
#SBATCH --gpus-per-node=4     # Specify total number of GPUs
#SBATCH -t 2:00:00            # Specify maximum time limit (hour: minute: second)
#SBATCH -A your_project       # Specify project name
#SBATCH -J JOBNAME            # Specify job name

module load Miniforge3/25.3.0-3
conda activate your_environment

export HF_HOME=/project/your_project_directory/hf/misc
export HF_DATASETS_CACHE=/project/your_project_directory/hf/datasets
export TRANSFORMERS_CACHE=/project/your_project_directory/hf/models
export HF_DATASETS_OFFLINE=1
export TRANSFORMERS_OFFLINE=1

srun python script.py
```

Using cray-python module on LANTA

➤ Load the cray-python module

```
username@lanta:~> ml av cray-python
----- /opt/cray/pe/lmod/modulefiles/core -----
cray-python/3.9.13.1  cray-python/3.10.10  cray-python/3.11.7 (D)
```

Use "module spider" to find all possible modules and extensions.

Use "module keyword key1 key2 ..." to search for all possible modules matching any of the "keys".

```
username@lanta:~> ml cray-python/3.11.7
```

➤ Activate the virtual environment

```
username@lanta:~> source /path/to/virtual/environment/bin/activate
(env-name) username@lanta:~>
```

Note: For R language users, use the cray-R module instead.

➤ Deactivate the virtual environment

```
(env-name) username@lanta:~> deactivate  
username@lanta:~>
```

➤ Unload the cray-python module

```
username@lanta:~> module unload cray-python
```

➤ Create the virtual environment

```
username@lanta:~> python -m venv ./myenv
```

➤ Activate the virtual environment

```
username@lanta:~> source ./myenv/bin/activate  
(myenv) username@lanta:~>
```

➤ Removal of the virtual environment

```
username@lanta:~> rm -rf ./myenv
```

➤ Running on Compute node

```
#!/bin/bash
#SBATCH -p compute
#SBATCH -N 1 -c 128
#SBATCH --ntasks-per-node=1
#SBATCH -t 120:00:00
#SBATCH -A tb901xxx
#SBATCH -J JOBNAME

module load cray-python/3.11.7
source /home/username/myenv/bin/activate

python3 file.py
```

Specify partition [Compute/Memory/GPU]
Specify number of nodes and processors per task
Specify tasks per node
Specify maximum time limit (hour: minute: second)
Specify project name
Specify job name

Load the cray-python module
Activate your virtual environment

Run your program or executable code

Note: Full node: -c 128, Half node: -c 64, 1/4 node: -c 32

➤ Running on GPU node

```
#!/bin/bash
#SBATCH -p gpu
#SBATCH -N 1 -c 16
#SBATCH --gpus-per-node=4
#SBATCH --ntasks-per-node=4
#SBATCH -t 120:00:00
#SBATCH -A tb901xxx
#SBATCH -J JOBNAME

module load cray-python/3.11.7
source /home/username/myenv/bin/activate

python3 file.py
```

Specify partition [Compute/Memory/GPU]
Specify number of nodes and processors per task
Specify number of GPU per task
Specify tasks per node
Specify maximum time limit (hour: minute: second)
Specify project name
Specify job name

Load the crat-python module
Activate your virtual environment

Run your program or executable code

Note: 1 GPU card: --gpus-per-node=1, 2 GPU cards: --gpus-per-node=2, 4 GPU cards: --gpus-per-node=4

Using Apptainer module on LANTA

➤ Load the Apptainer module

```
username@lanta:~> ml av Apptainer
----- /opt/cray/pe/lmod/modulefiles/core -----
Apptainer/1.1.6
```

Use "module spider" to find all possible modules and extensions.

Use "module keyword key1 key2 ..." to search for all possible modules matching any of the "keys".

```
username@lanta:~> ml Apptainer/1.1.6
```

➤ Unload the cray-python module

```
username@lanta:~> module unload Apptainer
```

➤ Using the apptainer pull command

```
username@lanta:~> apptainer pull [pull options] [output file] <URI>
```

➤ Source

<https://hub.docker.com/>

<https://catalog.ngc.nvidia.com/containers>

➤ Examples

From Docker

```
username@lanta:~> apptainer pull tensorflow.sif docker://tensorflow/tensorflow:latest-gpu
```

From NVIDIA NGC Catalog

```
username@lanta:~> apptainer pull pytorch.sif docker://nvcr.io/nvidia/pytorch:24.05-py3
```

➤ Running on Compute node

```
#!/bin/bash
#SBATCH -p compute
#SBATCH -N 1 -c 128
#SBATCH --ntasks-per-node=1
#SBATCH -t 120:00:00
#SBATCH -A tb901xxx
#SBATCH -J JOBNAME

module load Apptainer/1.1.6

apptainer exec -B $PWD:$PWD file.sif python3 file.py
```

Specify partition [Compute/Memory/GPU]
Specify number of nodes and processors per task
Specify tasks per node
Specify maximum time limit (hour: minute: second)
Specify project name
Specify job name

Load the Apptainer module

Run your program

Note: Full node: -c 128, Half node: -c 64, ¼ node: -c 32

➤ Running on GPU node

```
#!/bin/bash
#SBATCH -p gpu
#SBATCH -N 1 -c 16
#SBATCH --gpus-per-node=4
#SBATCH --ntasks-per-node=4
#SBATCH -t 120:00:00
#SBATCH -A tb901xxx
#SBATCH -J JOBNAME

module load Apptainer/1.1.6

apptainer exec --nv -B $PWD:$PWD file.sif python3 file.py # Run your program
```

Specify partition [Compute/Memory/GPU]
Specify number of nodes and processors per task
Specify number of GPU per task
Specify tasks per node
Specify maximum time limit (hour: minute: second)
Specify project name
Specify job name

Load the Apptainer module

Note: 1 GPU card: --gpus-per-node=1, 2 GPU cards: --gpus-per-node=2, 4 GPU cards: --gpus-per-node=4



Contact us

www.thaisc.io

thaisc-support@nstda.or.th